

NOTE: The proofs here are intended to provide the rough idea, and not necessarily give the full details.

1: Recurrences, recurrences

(a) In the recursion tree, we have $k := \log_4 n$ levels, and in level i , we have 4^{i-1} terms that are $n/4^{i-1}$, along with 4^i terms $T(n/4^i)$ terms. Thus the contribution of each level to the sum is n , and the contribution of the *last* level due to the $T()$ terms is $4^k T(1) = O(n)$. This gives an overall bound of $O(n \log n)$.

(b) We consider the recursion tree as before. Now, in level i , we have 4^{i-1} terms that are 1, along with the $T()$ terms as above. The contribution to the last level due to the $T()$ terms is $O(n)$ once again. Thus, the overall bound is

$$\left(\sum_{i=1}^{\log_4 n} 4^{i-1} \right) + O(n) = O(n).$$

(c) $T(n) = T(n-1) + n = T(n-2) + (n-1) + n = \dots = T(1) + 2 + \dots + n = \frac{n(n+1)}{2} + T(1) - 1$. This is clearly $O(n^2)$.

(d) It is clear that $T(n) \geq \sqrt{n}$ and that $T(n) \leq O(n)$. The latter can easily be seen by induction. Now, suppose we guess that the answer is $T(n) = Cn^\alpha$, for some $\alpha \in [1/2, 1]$. Then, for an inductive proof to go through, we will need

$$C(1/2^\alpha + 1/3^\alpha) + n^{1/2-\alpha} \leq C.$$

What is the smallest α for which we can find a C such that the above holds? We can do a quick search. Turns out $\alpha = 0.79$ works.

(e) $T(n) = T(n^{1/2}) + 4 = T(n^{1/4}) + 8 = \dots = T(n^{1/2^r}) + 4r$, for any r . If we set r such that $1/2^r = 1/\log n$, then the $T(n^{1/2^r})$ term becomes $T(1)$. Thus we have $T(n) = O(\log \log n)$.

(f) In the first case, the i th level in the tree has 3^i terms that are $T(n/2^i)$, and 3^{i-1} terms that are $(n^2/4^{i-1})$. Thus the ‘work’ done at the i th level is $n^2(3/4)^{(i-1)}$. Overall the tree has $k = \log_2 n$ levels. At the last level, the $T(1)$ terms add up to $O(3^k) = O(n^{\log_2 3})$. Thus the overall bound we get is $T(n) = O(n^2)$ (as $\sum_i (3/4)^i = O(1)$, and $\log_2 3 < 2$).

In the second case, the dominant terms are the $T(1)$ terms in the last level. This gives $T(n) = O(n^{\log_2 3})$.

In the final case, the work at each level works out to precisely $n^{\log_2 3}$ (irrespective of the level). Thus the overall bound we get is $T(n) = O(n^{\log_2 3} \log n)$.

2: Sorting nearby numbers

Algorithm: We maintain an array B of size M , in which the j th element is the number of times the number $\min_i A[i] + j$ appears in the array A . This array is initialized to zero, and is then populated in time $O(n)$, by performing a sweep through A .

Next, we can do a sweep through $B[]$, and using the counts, output the sorted version of A .

Running time: The two sweeps are both linear in the size of the corresponding arrays. This gives the bound $O(n + M)$.

3: Selecting in a Union

First, note that the answer can only lie in $A[0, \dots, k-1]$ or $B[0, \dots, k-1]$ (i.e., the remaining elements in the two arrays are irrelevant). Thus we will suppose that $N = k$.

Next, note that if we can find an r s.t. $A[r-1]$ satisfies $B[k-r-1] \leq A[r-1] \leq B[k-r]$, then $A[r-1]$ is the k th smallest element in the union. With this in mind, let us do a “binary search” in A as follows: start with $r = k/2$, and check if the above happens. If not, either $A[r-1] < B[k-r-1]$ or $A[r-1] > B[k-r]$. In the former case, the k th smallest element must be to the right of $A[r]$, thus we set $r = 3k/4$ (doing floor/ceiling as in regular binary search). Meanwhile, in the latter case, the k th smallest element must be to the left of $A[r]$, and we then proceed, setting $r = k/4$.

This process continues, and we either find the k th smallest element, or we find two consecutive indices $r-1, r$ with the guarantee that the answer lies between them. In this case, the k th smallest element in the union is $B[k-r-1]$.

4: Closest pair of restaurants in Manhattan

(a) Have points at $(1, 0), (4, 0), (5, 0), (8, 0)$. Then $x = 4.5$ is a vertical line that partitions the points into two sets of size $n/2$. In this case, $d = 3$, while the right answer is 1.

(b) Proof of hint: tile the $d \times d$ square S with 6 squares, each of size $(d/2) \times (d/3)$. Now, if we had 7 points in S , there have to be 2 points in one of the tiles (pigeon hole principle), and thus these two points have a Manhattan distance at most $d/2 + d/3 < d$.

Once we have this, we can reason about the algorithm as follows. First, the min pairwise distance can occur either between two points on the same side of x , or between points on different sides. In the former case, d will be the answer. In the latter case, let us denote by a and b the points that achieve this min distance. It is clear that the x-coordinates of a, b lie in $[x-d, x+d]$ (else the distance between a, b is $> d$). Thus we can throw away all the other points outside the strip, as step 3 does. Now consider sorting the points by y coordinates. In this ordering, we may suppose without loss of generality, that a appears before b . If there are ≤ 12 points between a and b in the ordering, then $\|a - b\|$ will be considered as a candidate distance, we will find it. Now, can there be ≥ 13 points in the ordering ‘between’ a and b ? Suppose there are. First, note that the difference in y -coordinates of a, b is at most d (else the distance is $> d$). See fig below.

At least one of R_1 and R_2 contain ≥ 7 points. Thus there are points in either R_1 or R_2 that are $< d$ apart. This contradicts the fact that d is the shortest distance in the recursive calls in step (2).

(c) (Skipping description – pretty straightforward). We get the $O(n \log n)$ because of the sorting in step 4 of the algorithm. The rest of step 4 takes $O(n)$ time.

[Recurrence easy to solve]




Figure 1: Figure showing possible relative positions of a, b

5: Linear Time Median

(You can look up any standard reference, e.g., the Dasgupta textbook.)