

# Lecture: Pipelining Extensions

---

- Topics: control hazards, multi-cycle instructions, pipelining equations

# Summary

---

- For the 5-stage pipeline, bypassing can eliminate delays between the following example pairs of instructions:

add/sub            R1, R2, R3  
add/sub/lw/sw    R4, R1, R5

lw        R1, 8(R2)  
sw        R1, 4(R3)

- The following pairs of instructions will have intermediate stalls:

lw                    R1, 8(R2)  
add/sub/lw        R3, R1, R4        or    sw    R3, 8(R1)

fmul        F1, F2, F3  
fadd        F5, F1, F4

# Problem 8

---

- Consider this 8-stage pipeline (RR and RW take a full cycle)




- For the following pairs of instructions, how many stalls will the 2<sup>nd</sup> instruction experience (with and without bypassing)?
  - ADD R3  $\leftarrow$  R1+R2  
ADD R5  $\leftarrow$  R3+R4
  - LD R2  $\leftarrow$  [R1]  
ADD R4  $\leftarrow$  R2+R3
  - LD R2  $\leftarrow$  [R1]  
SD R3  $\rightarrow$  [R2]
  - LD R2  $\leftarrow$  [R1]  
SD R2  $\rightarrow$  [R3]

# Problem 8

---

- Consider this 8-stage pipeline (RR and RW take a full cycle)



- For the following pairs of instructions, how many stalls will the 2<sup>nd</sup> instruction experience (with and without bypassing)?

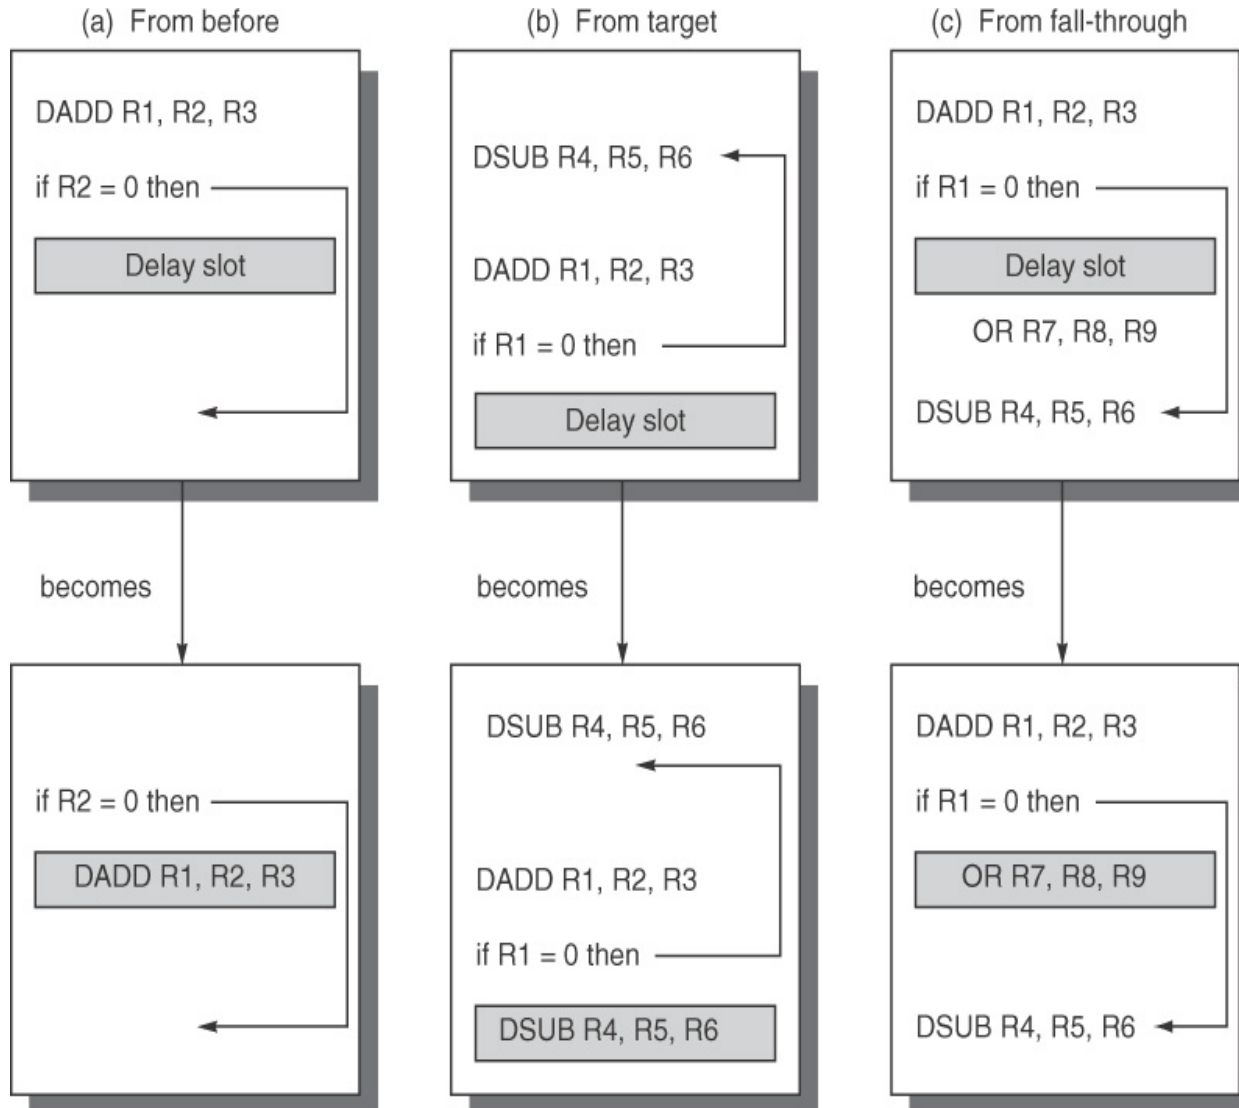
- |                             |                    |
|-----------------------------|--------------------|
| ▪ ADD R3 $\leftarrow$ R1+R2 |                    |
| ADD R5 $\leftarrow$ R3+R4   | without: 5 with: 1 |
| ▪ LD R2 $\leftarrow$ [R1]   |                    |
| ADD R4 $\leftarrow$ R2+R3   | without: 5 with: 3 |
| ▪ LD R2 $\leftarrow$ [R1]   |                    |
| SD R3 $\rightarrow$ [R2]    | without: 5 with: 3 |
| ▪ LD R2 $\leftarrow$ [R1]   |                    |
| SD R2 $\rightarrow$ [R3]    | without: 5 with: 1 |

# Control Hazards

---

- Simple techniques to handle control hazard stalls:
  - for every branch, introduce a stall cycle (note: every 6<sup>th</sup> instruction is a branch on average!)
  - assume the branch is not taken and start fetching the next instruction – if the branch is taken, need hardware to cancel the effect of the wrong-path instructions
  - predict the next PC and fetch that instr – if the prediction is wrong, cancel the effect of the wrong-path instructions
  - fetch the next instruction (branch delay slot) and execute it anyway – if the instruction turns out to be on the correct path, useful work was done – if the instruction turns out to be on the wrong path, hopefully program state is not lost

# Branch Delay Slots



# Problem 1

---

- Consider a branch that is taken 80% of the time. On average, how many stalls are introduced for this branch for each approach below:
  - Stall fetch until branch outcome is known
  - Assume not-taken and squash if the branch is taken
  - Assume a branch delay slot
    - You can't find anything to put in the delay slot
    - An instr before the branch is put in the delay slot
    - An instr from the taken side is put in the delay slot
    - An instr from the not-taken side is put in the slot

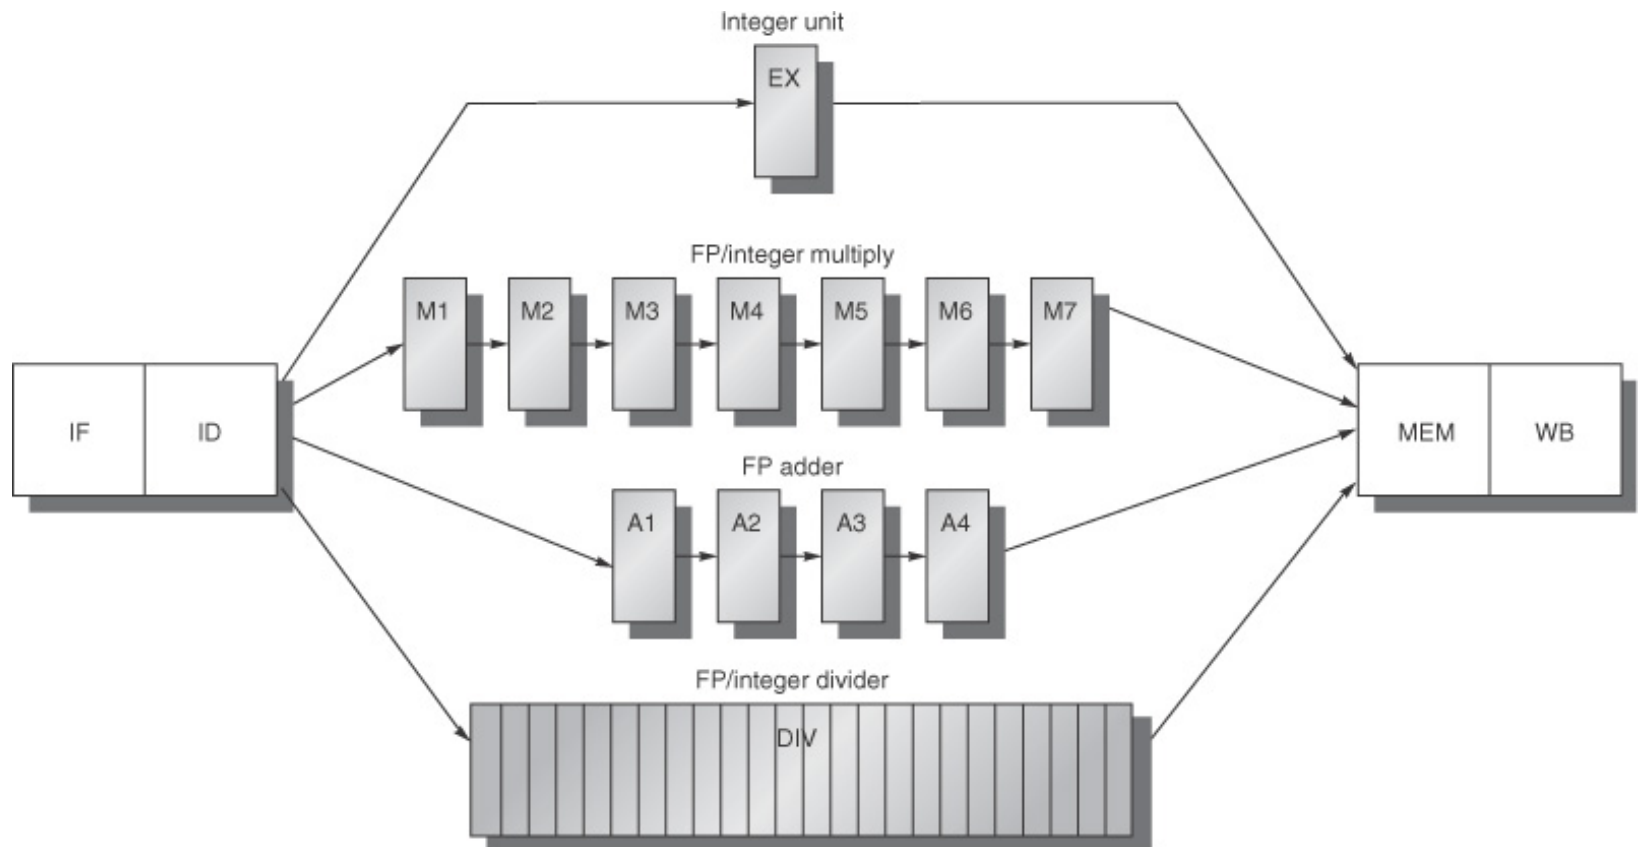
# Problem 1

---

- Consider a branch that is taken 80% of the time. On average, how many stalls are introduced for this branch for each approach below:
  - Stall fetch until branch outcome is known – 1
  - Assume not-taken and squash if the branch is taken – 0.8
  - Assume a branch delay slot
    - You can't find anything to put in the delay slot – 1
    - An instr before the branch is put in the delay slot – 0
    - An instr from the taken side is put in the slot – 0.2
    - An instr from the not-taken side is put in the slot – 0.8



# Multicycle Instructions



© 2007 Elsevier, Inc. All rights reserved.

# Effects of Multicycle Instructions

---

- Potentially multiple writes to the register file in a cycle
- Frequent RAW hazards
- WAW hazards (WAR hazards not possible)
- Imprecise exceptions because of o-o-o instr completion

Note: Can also increase the “width” of the processor: handle multiple instructions at the same time: for example, fetch two instructions, read registers for both, execute both, etc.

# Precise Exceptions

---

- On an exception:
  - must save PC of instruction where program must resume
  - all instructions after that PC that might be in the pipeline must be converted to NOPs (other instructions continue to execute and may raise exceptions of their own)
  - temporary program state not in memory (in other words, registers) has to be stored in memory
  - potential problems if a later instruction has already modified memory or registers
- A processor that fulfils all the above conditions is said to provide precise exceptions (useful for debugging and of course, correctness)

# Dealing with these Effects

---

- Multiple writes to the register file: increase the number of ports, stall one of the writers during ID, stall one of the writers during WB (the stall will propagate)
- WAW hazards: detect the hazard during ID and stall the later instruction
- Imprecise exceptions: buffer the results if they complete early or save more pipeline state so that you can return to exactly the same state that you left at


# Slowdowns from Stalls

---

- Perfect pipelining with no hazards  $\rightarrow$  an instruction completes every cycle (total cycles  $\sim$  num instructions)  
 $\rightarrow$  speedup = increase in clock speed = num pipeline stages
- With hazards and stalls, some cycles (= stall time) go by during which no instruction completes, and then the stalled instruction completes
- Total cycles = number of instructions + stall cycles
- Slowdown because of stalls =  $1 / (1 + \text{stall cycles per instr})$


# Pipelining Limits

---



Gap between indep instrs:  $T + T_{ovh}$

Gap between dep instrs:  $T + T_{ovh}$




Gap between indep instrs:

$$T/3 + T_{ovh}$$

Gap between dep instrs:

$$T + 3T_{ovh}$$



Gap between indep instrs:

$$T/6 + T_{ovh}$$

Gap between dep instrs:

$$T + 6T_{ovh}$$

Assume that there is a dependence where the final result of the first instruction is required before starting the second instruction

## Problem 2

---

- Assume an unpipelined processor where it takes 5ns to go through the circuits and 0.1ns for the latch overhead. What is the throughput for 20-stage and 40-stage pipelines? Assume that the P.O.P and P.O.C in the unpipelined processor are separated by 2ns. Assume that half the instructions do not introduce a data hazard and half the instructions depend on their preceding instruction.

## Problem 2

---

- Assume an unpipelined processor where it takes 5ns to go through the circuits and 0.1ns for the latch overhead. What is the throughput for 1-stage, 20-stage and 50-stage pipelines? Assume that the P.O.P and P.O.C in the unpipelined processor are separated by 2ns. Assume that half the instructions do not introduce a data hazard and half the instructions depend on their preceding instruction.
- 1-stage: 1 instr every 5.1ns
- 20-stage: first instr takes 0.35ns, the second takes 2.8ns
- 50-stage: first instr takes 0.2ns, the second takes 4ns
- Throughputs: 0.20 BIPS, 0.63 BIPS, and 0.48 BIPS



# ILP

---

- Instruction-level parallelism: overlap among instructions: pipelining or multiple instruction execution
- What determines the degree of ILP?
  - dependences: property of the program
  - hazards: property of the pipeline

# Static vs Dynamic Scheduling

---

- Arguments against dynamic scheduling:
  - requires complex structures to identify independent instructions (scoreboards, issue queue)
    - high power consumption
    - low clock speed
    - high design and verification effort
  - the compiler can “easily” compute instruction latencies and dependences – complex software is always preferred to complex hardware (?)

# Title

---

- Bullet