

# Lecture: SMT, Cache Hierarchies

---

- Topics: memory dependence wrap-up, SMT processors, cache access basics (Sections B.1-B.3, 2.1)

## Problem 3

---

- Consider the following LSQ and when operands are available. Estimate when the address calculation and memory accesses happen for each ld/st. Assume no memory dependence prediction.

	Ad. Op	St. Op	Ad.Val	Ad.Cal	Mem.Acc
LD R1 $\leftarrow$ [R2]	3		abcd		
LD R3 $\leftarrow$ [R4]	6		adde		
ST R5 $\rightarrow$ [R6]	5	7	abba		
LD R7 $\leftarrow$ [R8]	2		abce		
ST R9 $\rightarrow$ [R10]	1	4	abba		
LD R11 $\leftarrow$ [R12]	2		abba		

## Problem 3

---

- Consider the following LSQ and when operands are available. Estimate when the address calculation and memory accesses happen for each ld/st. Assume no memory dependence prediction.

	Ad. Op	St. Op	Ad.Val	Ad.Cal	Mem.Acc
LD R1 $\leftarrow$ [R2]	3		abcd	4	5
LD R3 $\leftarrow$ [R4]	6		adde	7	8
ST R5 $\rightarrow$ [R6]	5	7	abba	6	commit
LD R7 $\leftarrow$ [R8]	2		abce	3	7
ST R9 $\rightarrow$ [R10]	1	4	abba	2	commit
LD R11 $\leftarrow$ [R12]	2		abba	3	5

## Problem 4

---

- Consider the following LSQ and when operands are available. Estimate when the address calculation and memory accesses happen for each ld/st. Assume memory dependence prediction, with a default prediction that there is no dependence.

		Ad. Op	St. Op	Ad.Val	Ad.Cal	Mem.Acc
LD	R1 $\leftarrow$ [R2]	3		abcd		
LD	R3 $\leftarrow$ [R4]	6		adde		
ST	R5 $\rightarrow$ [R6]	4	7	abba		
LD	R7 $\leftarrow$ [R8]	2		abce		
ST	R9 $\rightarrow$ [R10]	8	3	abba		
LD	R11 $\leftarrow$ [R12]	1		abba		

## Problem 4

---

- Consider the following LSQ and when operands are available. Estimate when the address calculation and memory accesses happen for each ld/st. Assume memory dependence prediction, with a default prediction that there is no dependence.

		Ad. Op	St. Op	Ad.Val	Ad.Cal	Mem.Acc
LD	R1 $\leftarrow$ [R2]	3		abcd	4	5
LD	R3 $\leftarrow$ [R4]	6		adde	7	8
ST	R5 $\rightarrow$ [R6]	4	7	abba	5	commit
LD	R7 $\leftarrow$ [R8]	2		abce	3	4
ST	R9 $\rightarrow$ [R10]	8	3	abba	9	commit
LD	R11 $\leftarrow$ [R12]	1		abba	2	3/10

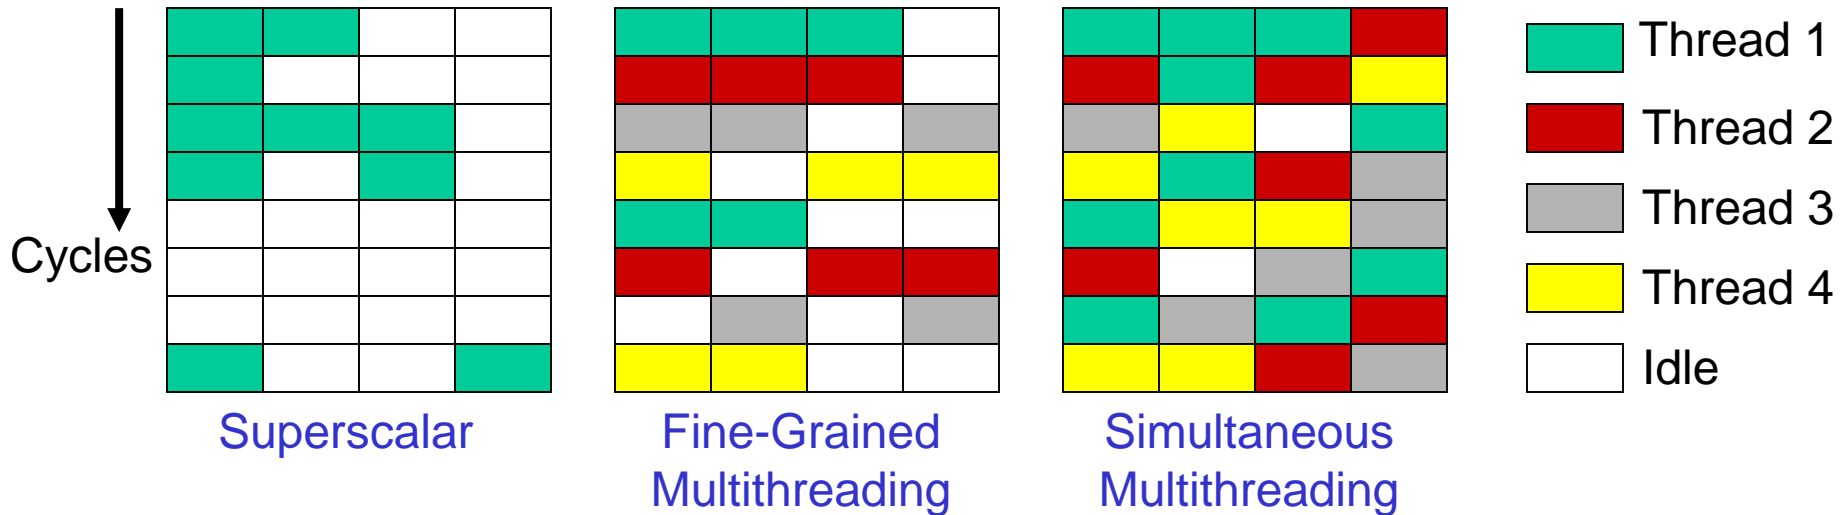
# Thread-Level Parallelism

---

- Motivation:
  - a single thread leaves a processor under-utilized for most of the time
  - by doubling processor area, single thread performance barely improves
- Strategies for thread-level parallelism:
  - multiple threads share the same large processor → reduces under-utilization, efficient resource allocation  
Simultaneous Multi-Threading (SMT)
  - each thread executes on its own mini processor → simple design, low interference between threads  
Chip Multi-Processing (CMP) or multi-core

# How are Resources Shared?

Each box represents an issue slot for a functional unit. Peak throughput is 4 IPC.



- Superscalar processor has high under-utilization – not enough work every cycle, especially when there is a cache miss
- Fine-grained multithreading can only issue instructions from a single thread in a cycle – can not find max work every cycle, but cache misses can be tolerated
- Simultaneous multithreading can issue instructions from any thread every cycle – has the highest probability of finding work for every issue slot


# What Resources are Shared?

---

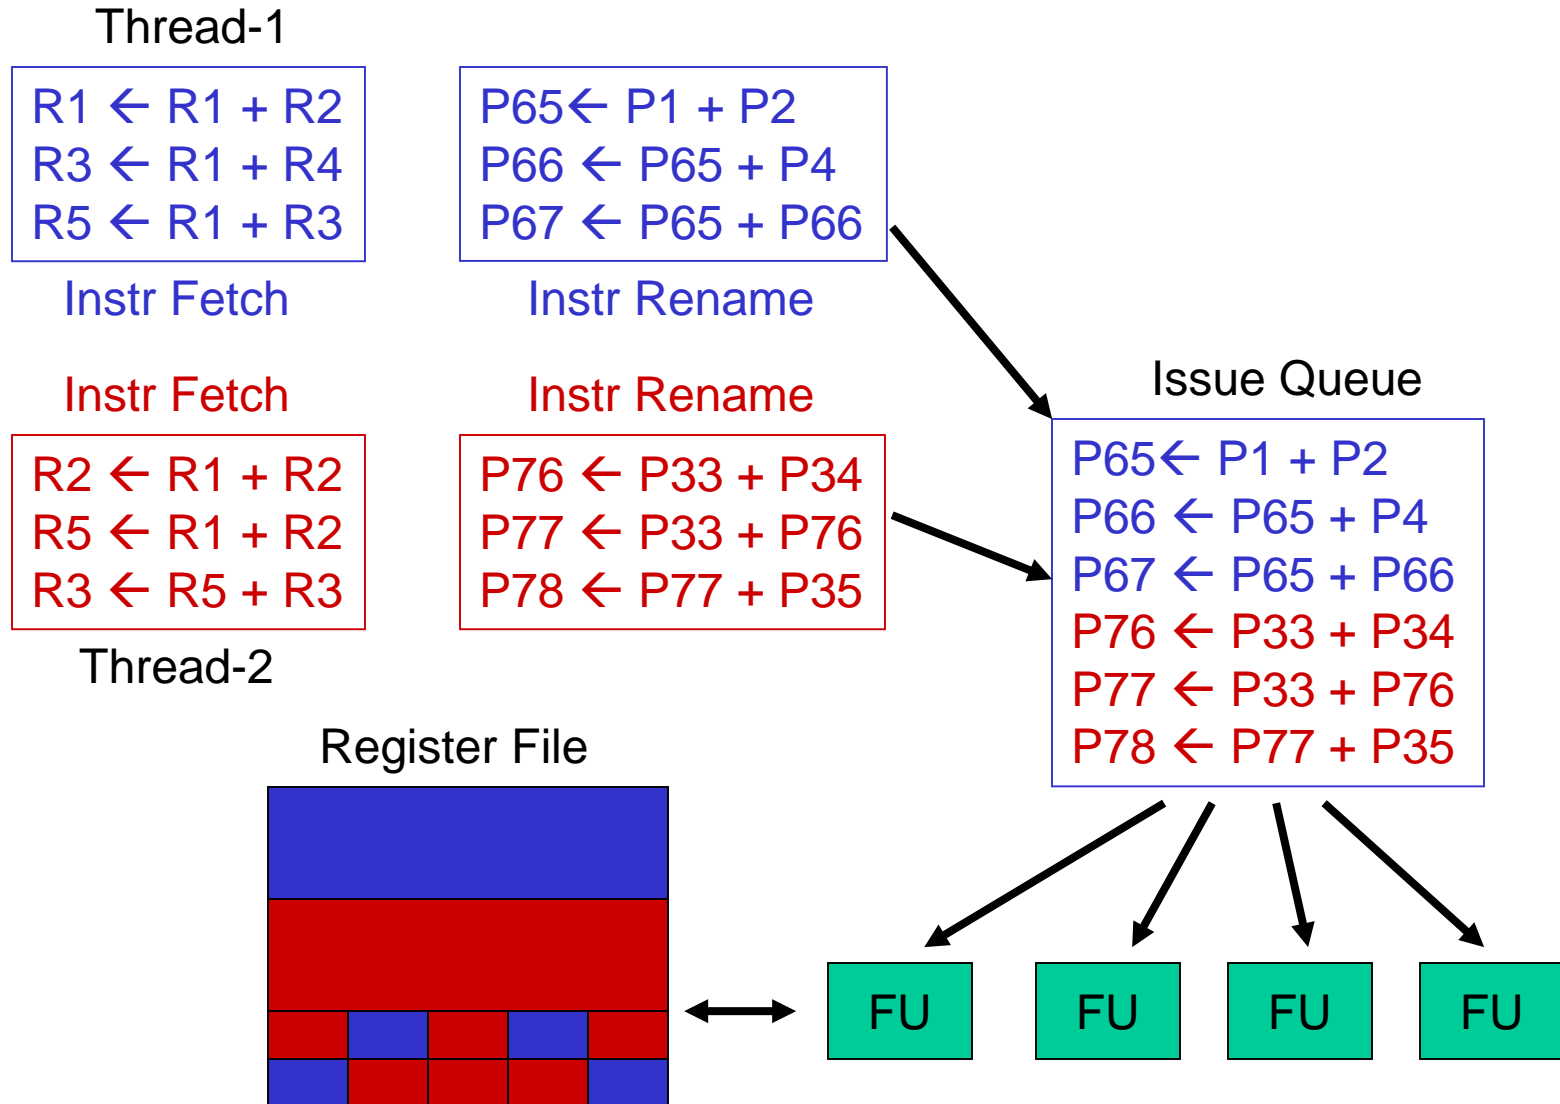
- Multiple threads are simultaneously active (in other words, a new thread can start without a context switch)
- For correctness, each thread needs its own PC, IFQ, logical regs (and its own mappings from logical to phys regs)
- For performance, each thread could have its own ROB/LSQ (so that a stall in one thread does not stall commit in other threads), I-cache, branch predictor, D-cache, etc. (for low interference), although note that more sharing → better utilization of resources
- Each additional thread costs a PC, IFQ, rename tables, and ROB – cheap!



# Pipeline Structure



# Resource Sharing



# Performance Implications of SMT

---

- Single thread performance is likely to go down (caches, branch predictors, registers, etc. are shared) – this effect can be mitigated by trying to prioritize one thread
- While fetching instructions, thread priority can dramatically influence total throughput – a widely accepted heuristic (ICOUNT): fetch such that each thread has an equal share of processor resources
- With eight threads in a processor with many resources, SMT yields throughput improvements of roughly 2-4


# Multi-Programmed Speedup

Benchmark	Best Speedup	Worst Speedup	Avg Speedup
gzip	1.48	1.14	1.24
vpr	1.43	1.04	1.17
gcc	1.44	1.00	1.11
mcf	1.57	1.01	1.21
crafty	1.40	0.99	1.17
parser	1.44	1.09	1.18
eon	1.42	1.07	1.25
perlbnk	1.40	1.07	1.20
gap	1.43	1.17	1.25
vortex	1.41	1.01	1.13
bzip2	1.47	1.15	1.24
twolf	1.48	1.02	1.16
wupwise	1.33	1.12	1.24
swim	1.58	0.90	1.13
mgrid	1.28	0.94	1.10
applu	1.37	1.02	1.16
mesa	1.39	1.11	1.22
galgel	1.47	1.05	1.25
art	1.55	0.90	1.13
equake	1.48	1.02	1.21
facerec	1.39	1.16	1.25
ammp	1.40	1.09	1.21
lucas	1.36	0.97	1.13
fma3d	1.34	1.13	1.20
sixtrack	1.58	1.28	1.42
apsi	1.40	1.14	1.23
Overall	1.58	0.90	1.20

- sixtrack and eon do not degrade their partners (small working sets?)
- swim and art degrade their partners (cache contention?)
- Best combination: swim & sixtrack  
worst combination: swim & art
- Static partitioning ensures low interference – worst slowdown is 0.9

# The Cache Hierarchy

---



# Problem 1

---

- Memory access time: Assume a program that has cache access times of 1-cyc (L1), 10-cyc (L2), 30-cyc (L3), and 300-cyc (memory), and MPKIs of 20 (L1), 10 (L2), and 5 (L3). Should you get rid of the L3?

# Problem 1


---

- Memory access time: Assume a program that has cache access times of 1-cyc (L1), 10-cyc (L2), 30-cyc (L3), and 300-cyc (memory), and MPKIs of 20 (L1), 10 (L2), and 5 (L3). Should you get rid of the L3?

With L3:  $1000 + 10 \times 20 + 30 \times 10 + 300 \times 5 = 3000$


Without L3:  $1000 + 10 \times 20 + 10 \times 300 = 4200$

# Accessing the Cache






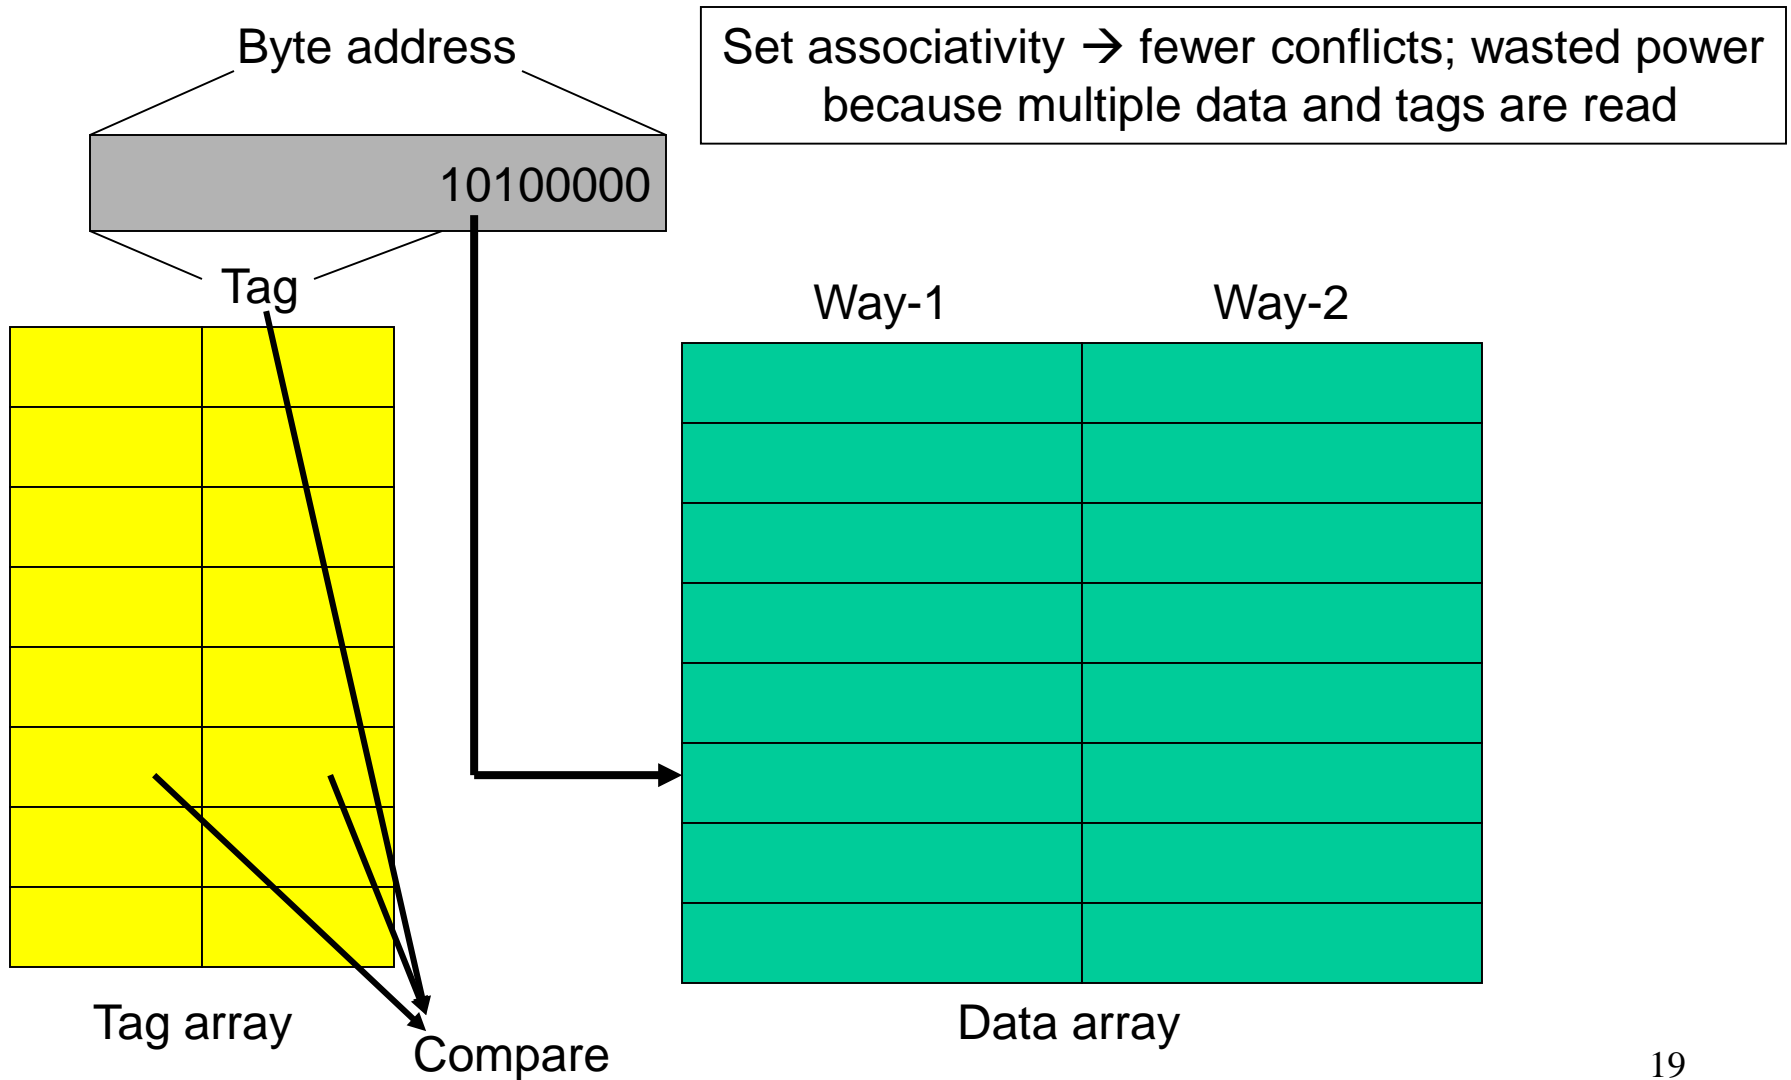
# The Tag Array



# Increasing Line Size



# Associativity



## Problem 2

---

- Assume a direct-mapped cache with just 4 sets. Assume that block A maps to set 0, B to 1, C to 2, D to 3, E to 0, and so on. For the following access pattern, estimate the hits and misses:

A B B E C C A D B F A E G C G A

## Problem 2

---

- Assume a direct-mapped cache with just 4 sets. Assume that block A maps to set 0, B to 1, C to 2, D to 3, E to 0, and so on. For the following access pattern, estimate the hits and misses:

A B B E C C A D B F A E G C G A  
M M H M M H M M H M M M M M

# Title

---

- Bullet